



Lecture 3: Loss Functions, Optimization, and Neural Network Overview

Areej Alasiry

CIS 6217 – Computer Vision for Data Representation

College of Computer Science, King Khalid University



Outline

1. Why Loss Functions?
2. Regression Losses (L1, L2)
3. Classification Losses (Hinge, Cross-Entropy)
4. Specialized Losses in CV
5. Gradient Descent Overview
6. Batch vs Stochastic vs Mini-batch
7. Optimization Improvements (Momentum, LR schedules)
8. Neuron Model & Activation Functions
9. Feed-Forward Neural Network Structure
10. Backpropagation
11. Lab Activity (Implement a small NN)
12. Summary

● Learning Outcomes

- Explain the role of loss functions in training machine learning models.
- Compare different loss functions used in computer vision tasks.
- Understand optimization techniques such as gradient descent and its variants.
- Describe the structure and operation of a basic neural network.
- Implement a simple feed-forward neural network and train it with backpropagation.

Loss Function

Measure the discrepancies between the prediction and the truth to guide training

- Why Loss Functions and Optimization



Regression Loss



- Mean Absolute Error (MAE)

$$L1 = \sum_{i=1}^n |y_i - y_i^p|$$

- Mean Squared Error (MSE)

$$L2 = \sum_{i=1}^n (y_i - y_i^p)^2$$


Truong, P. (2019) *Loss functions: Why, what, where or when?*, Medium. Available at: <https://phuctrt.medium.com/loss-functions-why-what-where-or-when-189815343d3f> (Accessed: 14 September 2025).



● Regression Loss Cont.

- Huber or Smooth Mean Absolute Error

$$L_{\delta}(y, f(x)) = \begin{cases} \frac{1}{2}(y - f(x))^2 & \text{for } |y - f(x)| \leq \delta \\ \delta|y - f(x)| - \frac{1}{2}\delta^2 & \text{otherwise.} \end{cases}$$


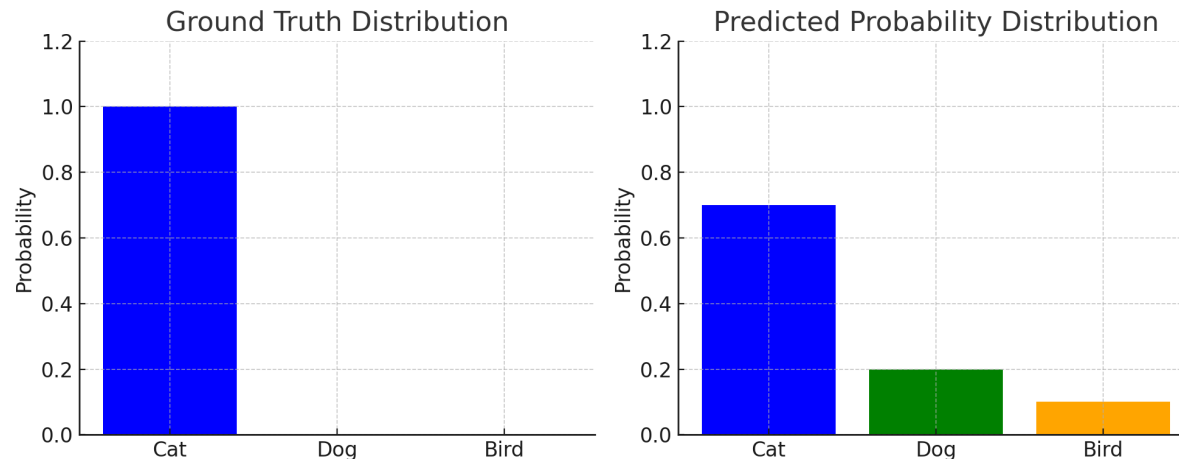


Truong, P. (2019) *Loss functions: Why, what, where or when?*, Medium. Available at: <https://phuctrt.medium.com/loss-functions-why-what-where-or-when-189815343d3f> (Accessed: 14 September 2025).

Classification Loss

- Cross-Entropy Loss (or Log Loss)

$$H(P, Q) = - \sum_i P(i) \log Q(i)$$



Truong, P. (2019) *Loss functions: Why, what, where or when?*, Medium. Available at: <https://phuctrt.medium.com/loss-functions-why-what-where-or-when-189815343d3f> (Accessed: 14 September 2025).



● Classification Loss

- Hinge Loss

$$J(\hat{y}, y) = \max(0, 1 - y_i \cdot \hat{y})$$

- Squared Hinge Loss

$$J(\hat{y}, y) = \sum_{i=0}^N (\max(0, 1 - y_i \cdot \hat{y}))^2$$




Truong, P. (2019) *Loss functions: Why, what, where or when?*, Medium. Available at: <https://phuctrt.medium.com/loss-functions-why-what-where-or-when-189815343d3f> (Accessed: 14 September 2025).

Loss in Image Segmentation

- Dice Loss = 1- Dice Coefficient
- Dice Coefficient (D)

$$D = \frac{2 \cdot |P \cap G|}{|P| + |G|}$$



Dice Loss Illustration

1 1 0 0

1 0 0 0

1 1 0 0

1 1 1 0

0 0 0 0

0 0 0 0


0 0 0 0

0 0 0 0

G

P

le (Green=Overlap, Gray=Missed, R




Gradient Descent

Gradually and iteratively solve a problem of function minimization .

Gradient Descent in Computer Vision

- **Gradient Descent** is an **optimization algorithm** used to minimize a loss (or cost) function by iteratively adjusting the model's parameters (weights).

$$\theta := \theta - \eta \cdot \nabla \theta L(\theta)$$



1. Model

$$\hat{y}_i = wx_i + b$$

2. Loss Function (MSE)

$$L(w, b) = \frac{1}{n} \sum_1^n (y_i - (wx_i + b))^2$$

3. Derivatives

Derivative wrt **w**:

$$\frac{\partial L}{\partial w} = \frac{-2}{n} \sum_{i=1}^n x_i (y_i - (wx_i + b))$$

Derivative wrt **b**:

$$\frac{\partial L}{\partial b} = \frac{-2}{n} \sum_{i=1}^n (y_i - (wx_i + b))$$

4. Update Rules (Gradient Descent)

$$w := w - \eta \cdot \frac{\partial L}{\partial w}$$

$$b := b - \eta \cdot \frac{\partial L}{\partial b}$$

5. Walkthrough Example

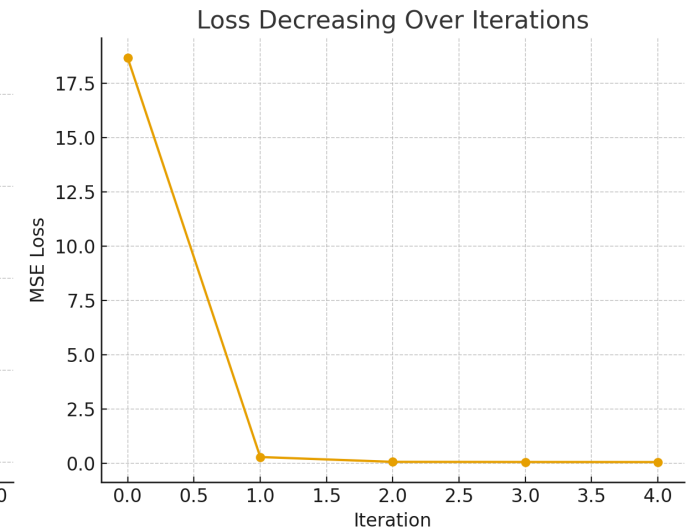
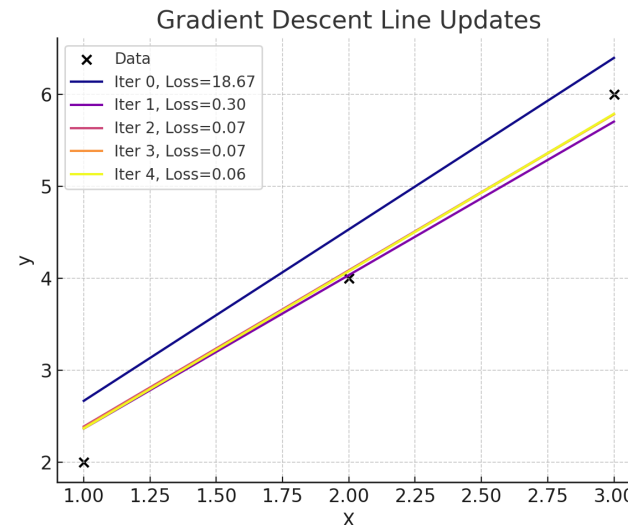
Dataset:

$$X = [1, 2, 3], y = [2, 4, 6]$$

Learning rate: $\eta = 0.1$.

Initial: $w = 0, b = 0$.

Iteration	w	b	Loss
0	1.866667	0.8	18.66667
1	1.671111	0.693333	0.296296
2	1.700741	0.686222	0.073376
3	1.70556	0.668681	0.067396
4	1.712898	0.652721	0.064165



● Variations of Gradient Descent

- Batch Gradient Descent

$$w := w - \eta \cdot \frac{1}{n} \sum_{i=1}^n \nabla_w L(x_i, y_i)$$

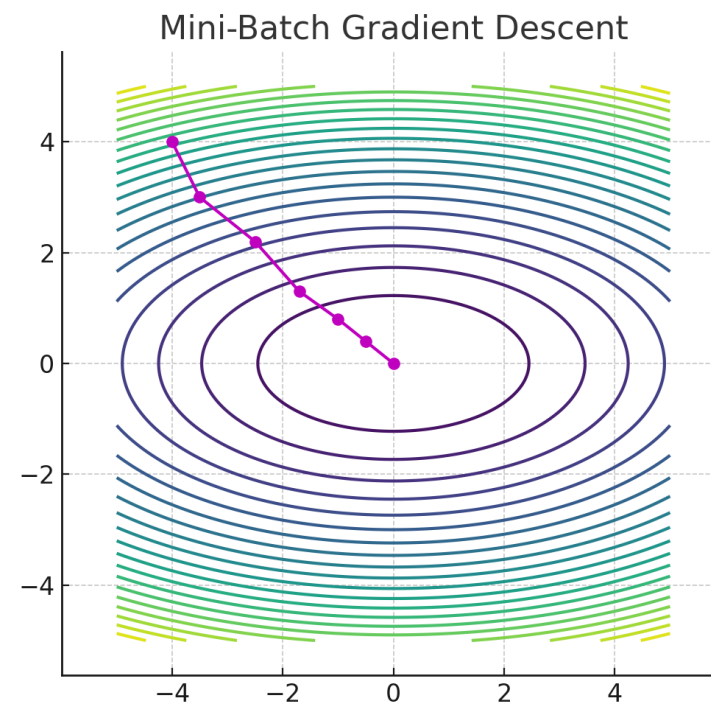
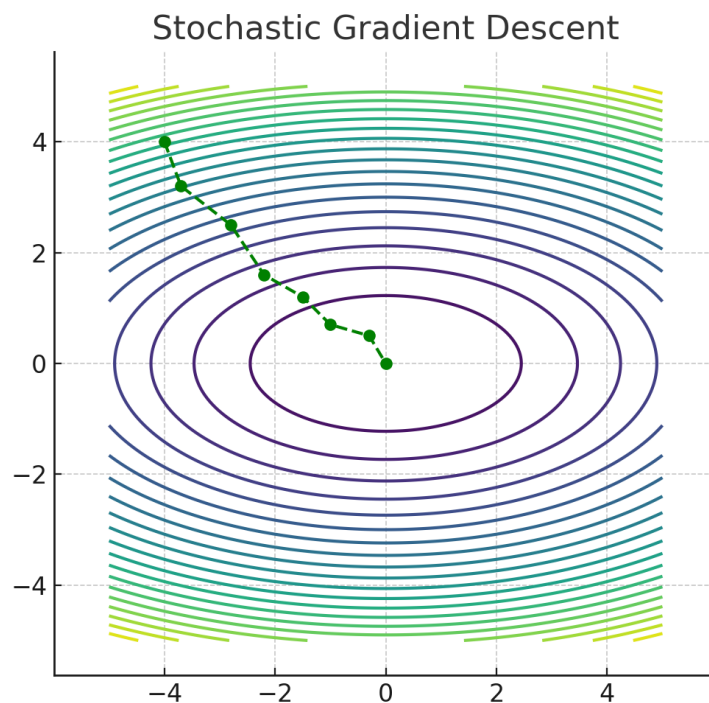

- Stochastic Gradient Descent

$$w := w - \eta \cdot \nabla_w L(x_i, y_i) \quad (\text{for a random sample } i)$$

- Mini Batch Gradient Descent

$$w := w - \eta \cdot \frac{1}{m} \sum_{i=1}^m \nabla_w L(x_i, y_i) \quad \text{where } m = \text{batch size}$$

● Illustration of Gradient Descent Variations




● Advanced Variations of Gradient Descent

- Momentum: Adds a **velocity term** to smooth updates.

$$v := \beta v + (1 - \beta) \nabla_w L, w := w - \eta v$$

- **RMSPprop**: Scales learning rate for each parameter based on past gradients.
- **Adam (Adaptive Moment Estimation)**: Combines **Momentum + RMSPprop**.

Neuron Model & Activation Function



● Neuron Model

Input features: x_1, x_2, \dots, x_n

Weights: w_1, w_2, \dots, w_n

Bias: b

- Weighted sum:


$$z = \sum_{i=1}^n w_i x_i + b$$

- Activation function:


$$a = f(z)$$

Activation Functions

1. Sigmoid




2. Tanh




3. ReLU (Rectified Linear Unit)

$$f(z) = \max(0, z)$$




4. Leaky ReLU

$$f(z) = \begin{cases} z & z \geq 0 \\ \alpha z & z < 0 \end{cases}$$



5. Softmax

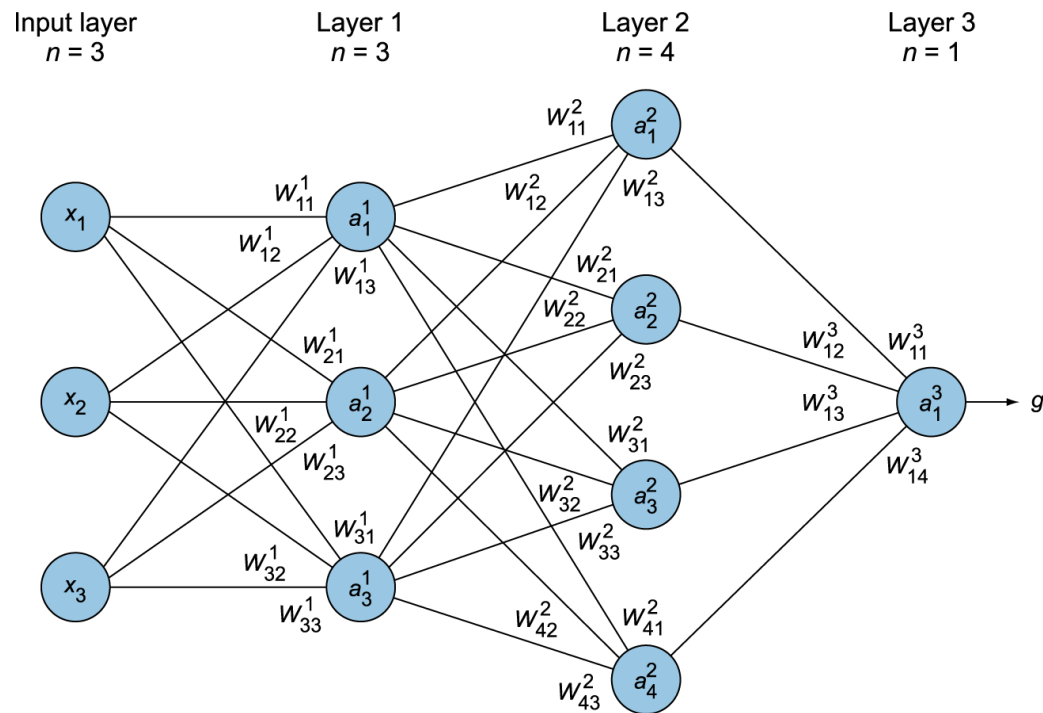
$$f(z_i) = \frac{e^{z_i}}{\sum_j e^{z_j}}$$



● Feed-Forward Neural Network

- A **feed-forward neural network** is the simplest type of artificial neural network, where information flows **in one direction** — from input → hidden layers → output, without cycles or feedback loops.
- **Structure**
 - **Input Layer:** Raw features (e.g., pixels of an image).
 - **Hidden Layers:** Neurons with weights, biases, and activations that transform input into abstract representations.
 - **Output Layer:** Produces final predictions (e.g., class probabilities with softmax).

Forward Pass



- **Computation Flow**

- Take inputs: x_1, x_2, \dots, x_n .
- Multiply by weights and add biases:

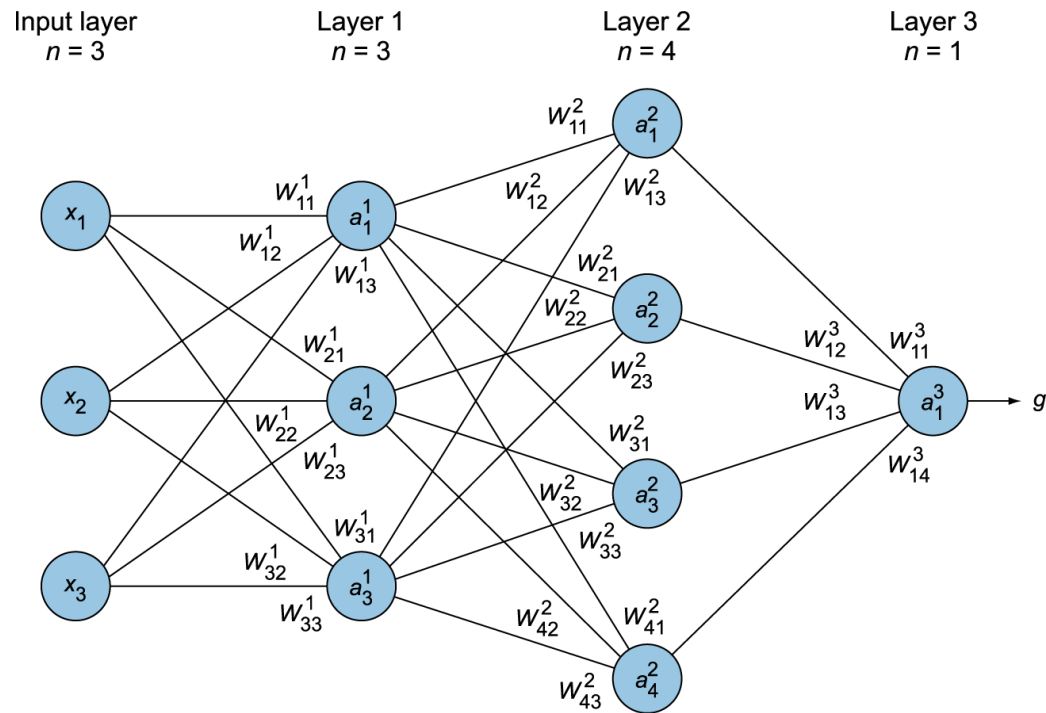
$$z^{(l)} = W^{(l)} a^{(l-1)} + b^{(l)}$$

- Apply activation function:

$$a^{(l)} = f(z^{(l)})$$

- Repeat through hidden layers until output.

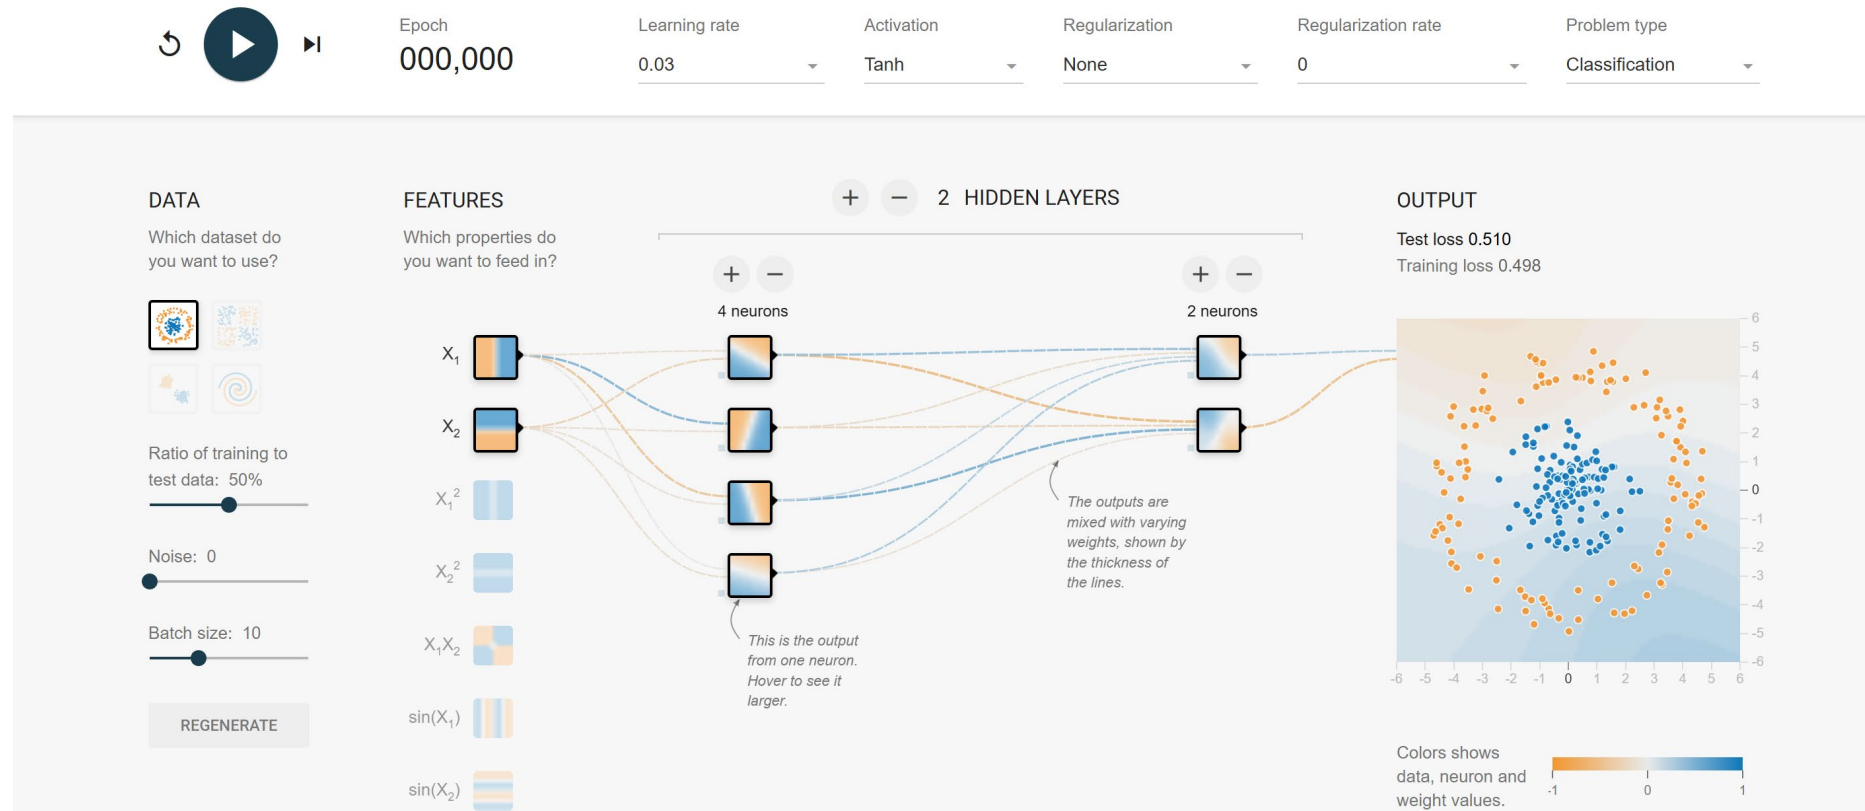
Calculation



$$\hat{y} = \sigma \left[\underbrace{\begin{bmatrix} w_{11}^3 & w_{12}^3 & w_{13}^3 & w_{14}^3 \end{bmatrix}}_{\text{Layer 3}} \cdot \sigma \left[\underbrace{\begin{bmatrix} w_{11}^2 & w_{12}^2 & w_{13}^2 \\ w_{21}^2 & w_{22}^2 & w_{23}^2 \\ w_{31}^2 & w_{32}^2 & w_{33}^2 \\ w_{41}^2 & w_{42}^2 & w_{43}^2 \end{bmatrix}}_{\text{Layer 2}} \cdot \sigma \left[\underbrace{\begin{bmatrix} w_{11}^1 & w_{12}^1 & w_{13}^1 \\ w_{21}^1 & w_{22}^1 & w_{23}^1 \\ w_{31}^1 & w_{32}^1 & w_{33}^1 \end{bmatrix}}_{\text{Layer 1}} \right] \right] \underbrace{\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}}_{\text{Input vector}}$$

Play around with NN

Tinker With a **Neural Network** Right Here in Your Browser.
Don't Worry, You Can't Break It. We Promise.



<https://tinyurl.com/2rekckkc>

Backpropagation

Forward Pass

Input flows through the network.

Compute predictions \hat{y} .

Compute loss $L(\hat{y}, y)$.


Backward Pass (Backpropagation)

Compute gradient of the loss wrt outputs of the last layer.

Apply **chain rule** layer by layer:

$$\frac{\partial L}{\partial w} = \frac{\partial L}{\partial a} \cdot \frac{\partial a}{\partial z} \cdot \frac{\partial z}{\partial w}$$

Update weights using gradient descent.



Ref: Deep Learning for Vision Systems by Mohamed Elgendy

● Summary

- Loss Functions: quantify error (L1/L2, Cross-Entropy, Dice).
- Optimization: Gradient Descent + its variants (Batch, SGD, Mini-Batch); advanced optimizers (Momentum, RMSProp, Adam).
- Neuron Model: weighted sum + bias + activation
- Feed-Forward NN: Input → Hidden Layers → Output, universal approximators.
- Backpropagation: chain rule-based algorithm to update weights and minimize loss.

● References

- Computer Vision: A Modern Approach – Forsyth & Ponce (2010)
- Extra: Deep Learning for Vision Systems by Mohamed Elgendy